

```

import gradio as gr
import torch
from transformers import AutoTokenizer, AutoModelForCausalLM
# -----
# Load Model & Tokenizer
# -----
model_name = "ibm-granite/granite-3.2-2b-instruct"
tokenizer = AutoTokenizer.from_pretrained(model_name)
model = AutoModelForCausalLM.from_pretrained(
    model_name,
    torch_dtype=torch.float16 if torch.cuda.is_available() else torch.float32,
    device_map="auto" if torch.cuda.is_available() else None,
)

if tokenizer.pad_token is None:
    tokenizer.pad_token = tokenizer.eos_token

# -----
# Generate AI Response
# -----
def generate_response(prompt, max_length=1024):
    inputs = tokenizer(prompt, return_tensors="pt", truncation=True, max_length=512)

    if torch.cuda.is_available():
        inputs = {k: v.to(model.device) for k, v in inputs.items()}

    with torch.no_grad():
        outputs = model.generate(
            **inputs,
            max_length=max_length,
            temperature=0.7,
            do_sample=True,
            pad_token_id=tokenizer.eos_token_id,
        )

    response = tokenizer.decode(outputs[0], skip_special_tokens=True)
    response = response.replace(prompt, "").strip()
    return response

# -----
# AI Functions
# -----
def city_analysis(city_name):
    prompt = f"""Provide a detailed analysis of {city_name} including:
1. Crime index and safety statistics
2. Accident rates and traffic safety
3. Overall safety assessment
"""
    return generate_response(prompt, max_length=1000)

def citizen_interaction(query):
    prompt = f"""You are a government assistant. Provide accurate and helpful information about:
Citizen Query: {query}
"""
    return generate_response(prompt, max_length=1000)

```

```

# -----
# Gradio Interface
# -----
with gr.Blocks() as app:
    gr.Markdown("# 🏙️ City Analysis & Citizen Services AI")

    with gr.Tabs():
        # Tab 1: City Analysis
        with gr.TabItem("City Analysis"):
            with gr.Row():
                with gr.Column():
                    city_input = gr.Textbox(
                        label="Enter City Name",
                        placeholder="e.g., New York, London, Mumbai...",
                    )
                    analyze_btn = gr.Button("Analyze City")
                with gr.Column():
                    city_output = gr.Textbox(
                        label="City Analysis (Crime Index & Accidents)", lines=15
                    )

            analyze_btn.click(city_analysis, inputs=city_input, outputs=city_output)

        # Tab 2: Citizen Services
        with gr.TabItem("Citizen Services"):
            with gr.Row():
                with gr.Column():
                    citizen_query = gr.Textbox(
                        label="Your Query",
                        placeholder="Ask about public services, government policies, civic issues..",
                        lines=4,
                    )
                    query_btn = gr.Button("Get Information")
                with gr.Column():
                    citizen_output = gr.Textbox(
                        label="Government Response", lines=15
                    )

            query_btn.click(
                citizen_interaction, inputs=citizen_query, outputs=citizen_output
            )

# -----
# Launch App
# -----
app.launch(share=True)

```

➔ /usr/local/lib/python3.12/dist-packages/huggingface\_hub/utils/\_auth.py:94: UserWarning:  
The secret `HF\_TOKEN` does not exist in your Colab secrets.  
To authenticate with the Hugging Face Hub, create a token in your settings tab (<https://huggingface.co/settings/tokens>).  
You will be able to reuse this secret in all of your notebooks.  
Please note that authentication is recommended but still optional to access public models or datasets.  
warnings.warn(

tokenizer\_config.json: 8.88k/? [00:00<00:00, 300kB/s]

vocab.json: 777k/? [00:00<00:00, 6.71MB/s]

merges.txt: 442k/? [00:00<00:00, 10.1MB/s]

tokenizer.json: 3.48M/? [00:00<00:00, 60.7MB/s]

added\_tokens.json: 100% 87.0/87.0 [00:00<00:00, 1.92kB/s]

special\_tokens\_map.json: 100% 701/701 [00:00<00:00, 31.5kB/s]

config.json: 100% 786/786 [00:00<00:00, 21.6kB/s]

`torch\_dtype` is deprecated! Use `dtype` instead!

model.safetensors.index.json: 29.8k/? [00:00<00:00, 2.98MB/s]

Fetching 2 files: 100% 2/2 [03:09<00:00, 189.49s/it]

model-00001-of-00002.safetensors: 100% 5.00G/5.00G [03:09<00:00, 90.1MB/s]

model-00002-of-00002.safetensors: 100% 67.1M/67.1M [00:05<00:00, 10.4MB/s]

Loading checkpoint shards: 100% 2/2 [00:19<00:00, 7.94s/it]

generation\_config.json: 100% 137/137 [00:00<00:00, 17.5kB/s]

Colab notebook detected. To show errors in colab notebook, set debug=True in launch()

\* Running on public URL: <https://a049d341910332e594.gradio.live>

This share link expires in 1 week. For free permanent hosting and GPU upgrades, run `gradio deploy`



## City Analysis & Citizen Services AI

City Analysis

Citizen Services

Enter City Name

e.g., New York, London, Mumbai...

Analyze City

City Analysis (Crime Index & Accidents)

